

CS2109S: Introduction to AI and Machine Learning

**Lecture 1:
Intro to CS2109S &
Artificial Intelligence**

12 Aug 2022







Rizki



Jonathan



Zitai




Rafeed



Kuluhan



Soo Han



Austen



Yuming



Jin Feng



Ming Chong



Aloysius



Rahul



Audrey




Arnav



Nicholas



Nyamdavaa



Gabriel

Rights Infringements on NUS Course Materials

All course participants (including permitted guest students) who have access to the course materials on LumiNUS or any approved platforms by NUS for delivery of NUS modules are not allowed to re-distribute the contents in any forms to third parties without the explicit consent from the module instructors or authorized NUS officials

Why are
you here?

Why are we here

- We were teaching the same things I taught in 2006/7 in CS3243
- Machine/Deep Learning is a big deal.
- CS2109S is now compulsory for all CS students
- Some CS majors might only take one AI class before graduation.

“old”
system

Intro to “classic” AI

CS3243

Machine learning

CS3244

“Too much Math”

CS2109S

Advanced ML

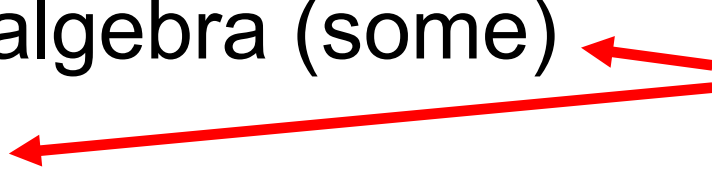
Advanced AI

Less Math (maybe)

More application (coding)

Focus on key ideas & intuitions

Course Pre-Requisites

- CS1101S, CS1010S or equivalent
 - CS1231 or equivalent
 - MA1521 (calculus)
 - CS2040S or equivalent
 - Linear algebra (some)
 - Python
- Problem Set 0
- language-agnostic
- 

Warning: we are
still fine tuning
this class.

Topics

1. Search (“classic AI”)
2. Decision Trees
3. Linear Regression
4. Gradient descent
5. Logistic Regression

Topics


- 6. Support Vector Machines
- 7. Artificial Neural networks
- 8. Unsupervised learning
- 9. AI & Ethics

Textbook

Russell and Norvig (2021)
Artificial Intelligence: A
Modern Approach (4th Ed)

AIMA

**No need
to buy!**



Don't memorize
anything

**Focus on concepts
& understanding!**

There will be
webcast

**But I don't control when
it appears on Luminus**

IMPORTANT

**Tutorial
scheduling survey**

Due 14 Aug 23:59!

Plagiarism Policy

- Okay to discuss, cite them
- Do not take away any code from your discussions.
- We will run plagiarism checker

Plagiarism Policy

- Do not use code from seniors
- Do not publish code on GitHub!!
- We will run plagiarism checker against assignments from previous batches

Learn
how to learn

Learn to
solve problems

Homework
needs to be
done to learn

Shouldn't be high stakes

Plagiarism
will be
a problem

Homework shouldn't be high
stakes

Students need
a fair grade

So how?

Grading

- Homework needs to be done to learn ⇒ Homework needs to be a high weightage
 - Plagiarism will be a problem ⇒ Homework cannot have a high weightage
 - Students need a fair grade
- Life is hard. 😞


Grading

- Homework needs to be done to learn ⇒ Homework needs to be a high weightage
- Plagiarism will be a problem ⇒ Homework easy to get full marks
- Students need a fair grade ⇒ Exams mainly determine final grade weightage

Grading

expected
to get 30%

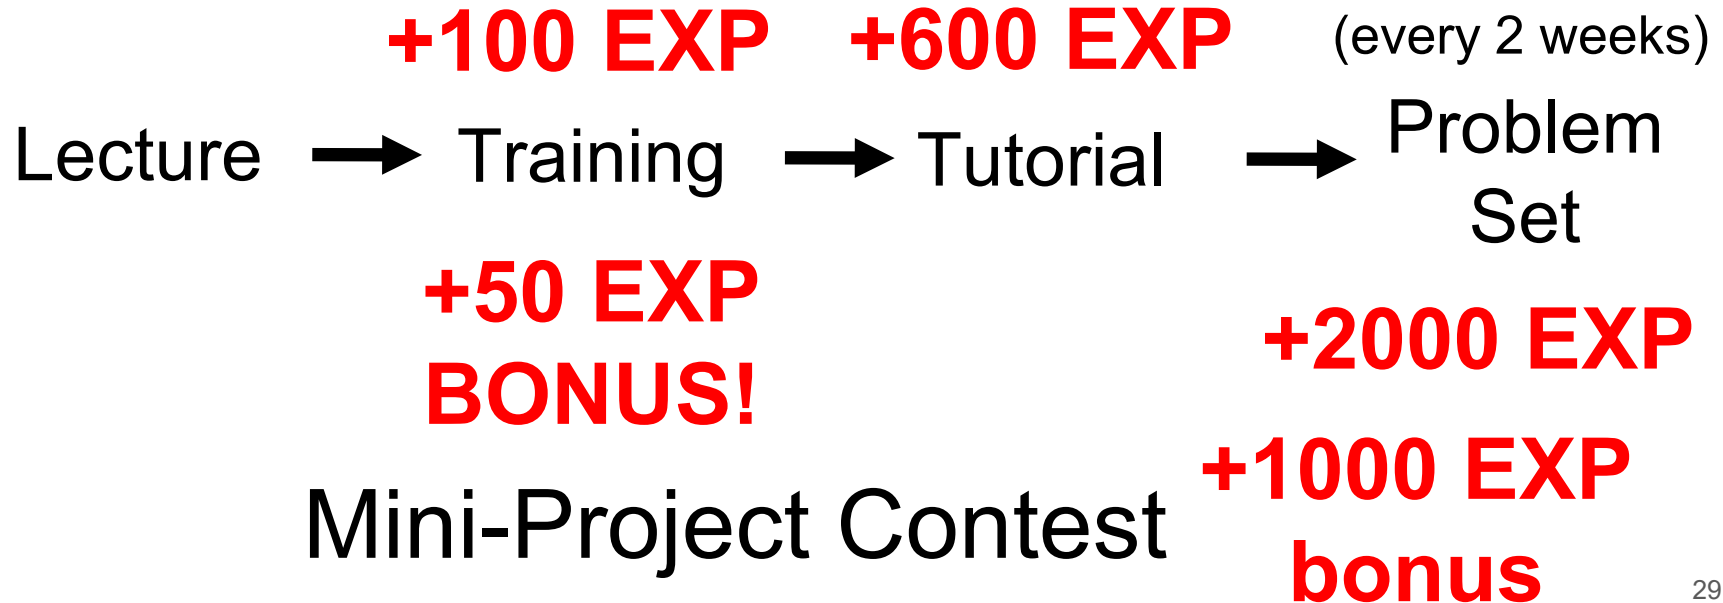
- Coursemology (CA) – 30%
- Mini-Project – 10%
- Midterm Assessment – 30%
- Final (Practical) Exam – 30%



Coursemology, not Luminus/Canvas

But will use Luminus
Gradebook to record marks!

Gamified CA (30%)



Please focus on
learning, not
grades

Gamified CA (30%)

Turn up for class

Coding Practice
(every 2 weeks)

Lecture → Training → Tutorial → Problem Set

Exploration

Mini-Project Contest

Enough EXP?

Level Up!

Final Level is CA grade

Background Survey

+100 EXP

Late Policy

- Up to 1 hour, no penalty
- Up to 24 hours, -20%
- Beyond 24 hours, -50%

Need extension, ask early

Assessment (Face-to-Face)

- Midterm Test (30%)
 - 30 Sep@10am (Fri, Week 7)
 - Venue: TBA **Might be evening because of venue**
- Final Exam (30%)
 - 25 Nov@2.30 pm (Friday)
 - Arranging for Practical Exam**

Exams

- Closed book, open-sheet
(as good as open book)
- Focus on application,
NOT memorization


Questions?

Who wants a C for CS2109S?

- **Don't procrastinate**
- Start on homework and projects early
- Students who wait till the last minute to do the assignments invariably fail to debug their programs properly


Workload is front-loaded

What is AI?





Credit: Wikipedia

Turing Test (1950)



CAPTCHA


(Completely Automated Public Turing test to tell Computers and Humans Apart)



Please click each image containing a seaplane



If there are None, click Skip



**“Acting
Humanly”**

What else
is considered
AI?

DeepBlue (1997)



Credit: IEEE Spectrum

AlphaGo (2015)



Credit: Guardian

iRobot Roomba (2002)



Credit: Wikipedia

Watson (2011) vs Jeopardy






Credit: NYTimes



Credit: Vulcan Post



Credit: Technology Review



Does it
matter? 😊

Learning
how to learn

Learning
to ~~solve problems~~
frame problems

Approach AI as Problem Solving

Computer Science
is the study of
abstractions

Managing Complexity

Functional Abstraction



Picture → **Function** → Word/number

Picture → **Function** → Chihuahua
/muffin

Email → **Function** → Spam
/not spam

Black
& white
photo → **Function** → Coloured
photo




How do we write this function?

Is this sufficient?



How would we model self-driving?

Intelligent Agents



Agents

- An agent is anything that can be viewed as **perceiving its environment through sensors** and **acting upon that environment through actuators**
- Human agent: eyes, ears, and other organs for sensors; hands, legs, mouth, and other body parts for actuators
- Robotic agent: cameras and infrared range finders for sensors; various motors for actuators

Agent design

The agent function maps from percept histories to actions:

$$[f: \mathcal{P}^* \rightarrow \mathcal{A}]$$

The agent program runs on the physical architecture to produce f

agent = architecture + program

How would an agent know
to do the right thing?

→ need a measure of
goodness

Performance measure

Performance Measure

1. Best for whom?
2. What are we optimising?
3. What information is available?
4. Unintended effects
5. What are the costs?

Performance vs. Cost

Rational agent

A rational agent will choose an action that is **expected to maximize its performance measure**, given the evidence provided by

- the **percept sequence**; and
- the **built-in knowledge** it has.

Rationality \neq Omniscience

Agents can perform actions in order to modify future percepts so as to obtain useful information (information gathering, exploration)

An agent is **autonomous** if its behavior is determined by its own experience (with ability to learn and adapt)

Defining the Problem: PEAS

1. **P**erformance measure
2. **E**nvironment
3. **A**ctuators
4. **S**ensors

Autonomous Driving

Performance Measure

1. Safety
2. Speed
3. Legal
4. Comfort

Environment

1. Roads
2. Weather/visibility
3. Other vehicles
4. Pedestrians/obstacles

Actuators

1. Steering wheel
2. Accelerator
3. Brake
4. Signal
5. Horn

Sensors

1. Camera
2. LIDAR
3. Speedometer
4. GPS
5. Engine sensors

Medical Diagnosis System

Performance Measure

1. Health outcome
2. Cost
3. Lawsuits

Environment

1. Patient
2. Hospital (equipment, etc.)
3. Staff

Actuators

Screen display
(questions, test,
diagnoses, treatments,
referrals)

Sensors

1. Keyboard
2. Medical
Readings
3. Medical History

Interactive Coding Tutor

Performance Measure

Maximize student's
score on test/Leetcode

Environment

1. Questions available
2. Language(s)

Actuators

Screen display
(questions, corrections,
suggestions)

Sensors

1. Keyboard
2. Input from
database

Characterizing the environment

Fully observable (vs. partially observable): An agent's sensors give it access to the complete state of the environment at each point in time.

Deterministic (vs. stochastic): The next state of the environment is completely determined by the current state and the action executed by the agent. (If the environment is deterministic except for the actions of other agents, then the environment is **strategic**)

Characterizing the environment

Episodic (vs. sequential): The agent's experience is divided into atomic "episodes" (each episode consists of the agent perceiving and then performing a single action), and the choice of action in each episode depends only on the episode itself.

"Memoryless"

Static (vs. dynamic): The environment is unchanged while an agent is deliberating. (The environment is **semi-dynamic** if the environment itself does not change with the passage of time but the agent's performance score does)

Characterizing the environment

Discrete (vs. continuous): A limited number of distinct, clearly defined percepts and actions.

Single agent (vs. multi-agent): An agent operating by itself in an environment.

Environment types

	Chess with a clock	Chess without a clock	Autonomous driving
Fully observable	Yes	Yes	No
Deterministic	Strategic	Strategic	No
Episodic	Yes	Yes	No
Static	Semi	Yes	No
Discrete	Yes	Yes	No
Single agent	No	No	No

The environment type largely determines the agent design
The real world is (of course) partially observable, stochastic, sequential, dynamic, continuous, multi-agent

Implementing Agents

An agent is completely specified by the agent function mapping percept sequences to actions

Goal: find a way to implement the rational agent function concisely

Playing Tic-Tac-Toe





Table-lookup agent

```
function TABLE-DRIVEN-AGENT(percept) returns action  
  static: percepts, a sequence, initially empty  
           table, a table of actions, indexed by percept sequences, fully specified  
  
  append percept to the end of percepts  
  action  $\leftarrow$  LOOKUP(percepts, table)  
  return action
```

Possible drawbacks:

- Storage
- No autonomy/cannot react to changes
- Might need a long time to compute/learn the table entries

Vacuum-cleaner world



Percepts: location and contents,
e.g., [A,Dirty]

Actions: *Left, Right, Suck, NoOp*


Vacuum Cleaner Agent

function REFLEX-VACUUM-AGENT([*location, status*]) **returns** an action

if *status* = *Dirty* **then return** *Suck*

else if *location* = *A* **then return** *Right*

else if *location* = *B* **then return** *Left*



Lookup table as case
statements

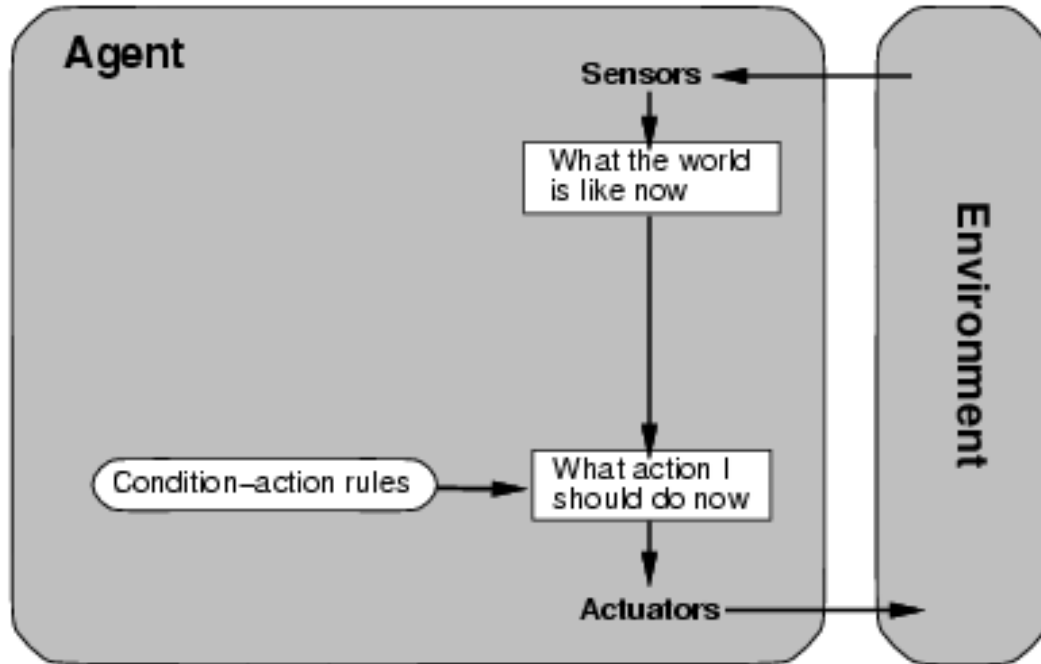
Models for Agent Organization

1. Simple reflex agents
2. Model-based reflex agents
3. Goal-based agents
4. Utility-based agents
5. Learning agents

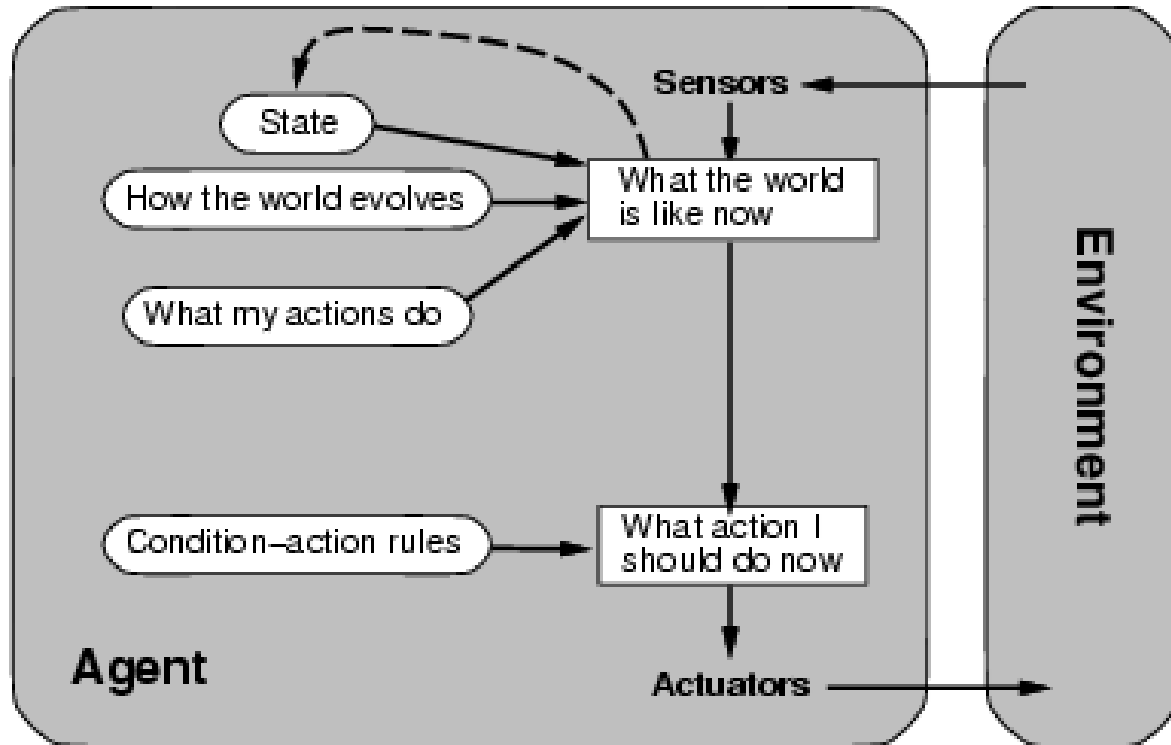


Increasing
complexity

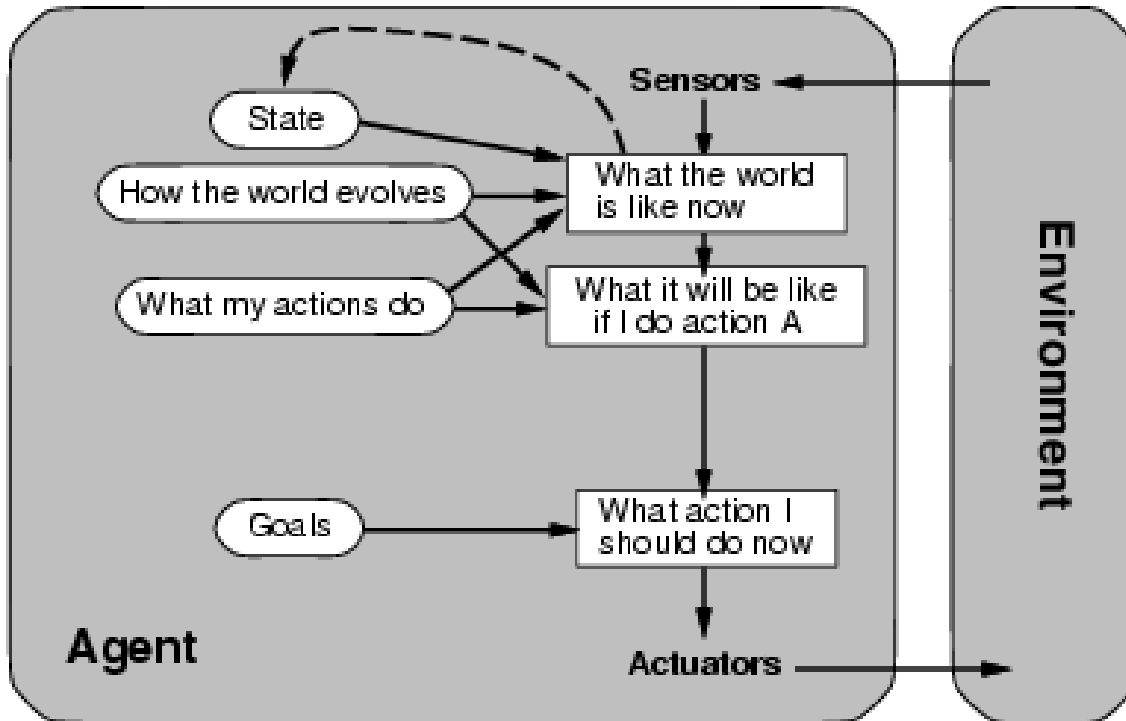
Simple Reflex Agent



Model-based Reflex Agent

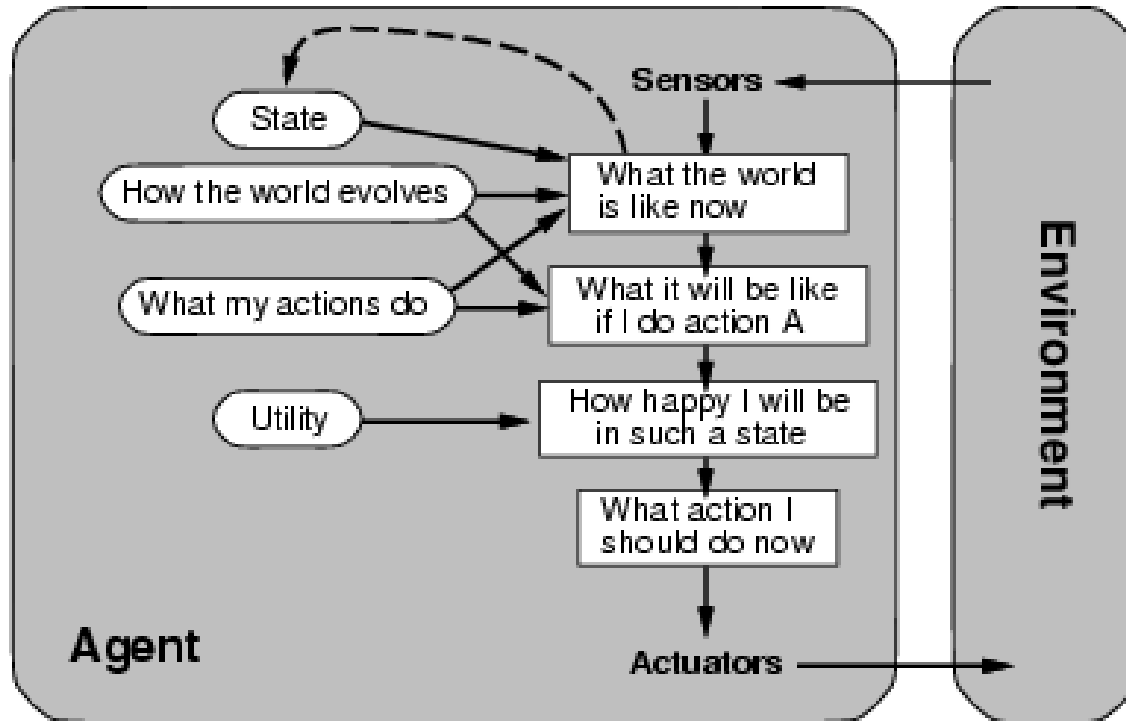


Goal-based Agent

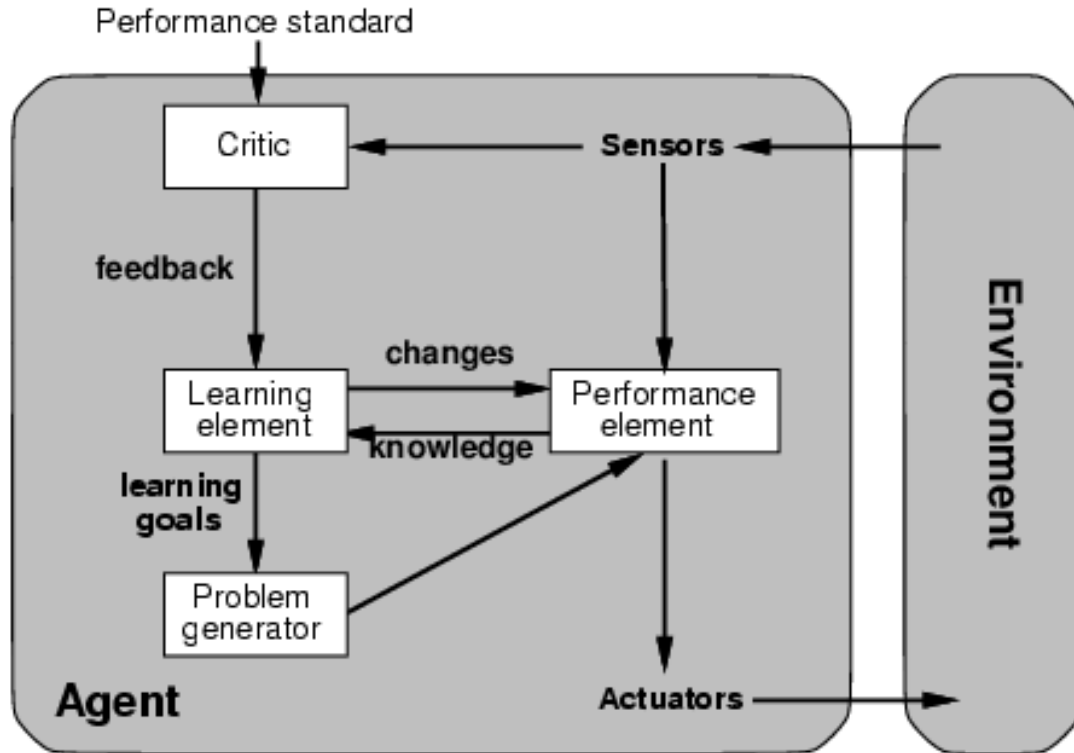


from
Economics

Utility-based Agent



Learning Agent



A learning agent might not
always maximize the
performance measure!

There are no model
answers in life.

Problem set solutions
will not be released

Exploitation vs Exploration

An agent operating in the real world must often choose between:

- maximizing its expected utility according to its current knowledge about the world; and
- trying to learn more about the world

since this may improve its future gains.

This problem is known as the trade-off between ***exploitation*** and ***exploration***

Summary

- Frameworks to think about and understand general AI problems
- PEAS framework
 - Performance Measure: defines “good” solution in solution space
 - Environment: determines what we can and cannot do
 - Actuators: output
 - Sensors: inputs
- Typical Solutions (different classes of agents)
 - Implementing a lookup-table
 - Reflex, Model-based, Goal-based, Utility-based, Learning

Summary

- Why do we care about these models?

Condition	Algorithm	Time Complexity
No Negative Weight Cycles	Bellman-Ford Algorithm	$O(VE)$
On Unweighted Graph (or equal weights)	BFS	$O(V + E)$
No Negative Weights	Dijkstra's Algorithm	$O((V + E)\log V)$
On Tree	BFS / DFS	$O(V)$
On DAG	Topological Sort	$O(V + E)$

2040S
anyone?

SSSP!

1. Tradeoff: Performance vs. Cost
2. Exploitation vs. Exploration

Do Lecture 1 Training!

Free +100 EXP

+50 EARLY BIRD
BONUS

Please start on
Problem Set 0
ASAP!

Due 20 Aug (Sat) 23:59